

R for Data science

```
library(nycflights13)
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.0.9000 --

## v ggplot2 3.3.0      v purrr 0.3.3
## v tibble 3.0.1       v dplyr 0.8.5
## v tidyr 1.0.2        v stringr 1.4.0
## v readr 1.3.1        v forcats 0.5.0

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()      masks stats::lag()

str(flights)

## tibble [336,776 x 19] (S3: tbl_df/tbl/data.frame)
## $ year      : int [1:336776] 2013 2013 2013 2013 2013 2013 2013 2013 2013 2013 ...
## $ month     : int [1:336776] 1 1 1 1 1 1 1 1 1 1 ...
## $ day       : int [1:336776] 1 1 1 1 1 1 1 1 1 1 ...
## $ dep_time  : int [1:336776] 517 533 542 544 554 554 555 557 557 558 ...
## $ sched_dep_time: int [1:336776] 515 529 540 545 600 558 600 600 600 600 ...
## $ dep_delay : num [1:336776] 2 4 2 -1 -6 -4 -5 -3 -3 -2 ...
## $ arr_time  : int [1:336776] 830 850 923 1004 812 740 913 709 838 753 ...
## $ sched_arr_time: int [1:336776] 819 830 850 1022 837 728 854 723 846 745 ...
## $ arr_delay : num [1:336776] 11 20 33 -18 -25 12 19 -14 -8 8 ...
## $ carrier   : chr [1:336776] "UA" "UA" "AA" "B6" ...
## $ flight    : int [1:336776] 1545 1714 1141 725 461 1696 507 5708 79 301 ...
## $ tailnum   : chr [1:336776] "N14228" "N24211" "N619AA" "N804JB" ...
## $ origin    : chr [1:336776] "EWR" "LGA" "JFK" "JFK" ...
## $ dest      : chr [1:336776] "IAH" "IAH" "MIA" "BQN" ...
## $ air_time  : num [1:336776] 227 227 160 183 116 150 158 53 140 138 ...
## $ distance  : num [1:336776] 1400 1416 1089 1576 762 ...
## $ hour      : num [1:336776] 5 5 5 5 6 5 6 6 6 6 ...
## $ minute    : num [1:336776] 15 29 40 45 0 58 0 0 0 0 ...
## $ time_hour : POSIXct[1:336776], format: "2013-01-01 05:00:00" "2013-01-01 05:00:00" ...

# filter() are combined with "and": every expression must be true in order
# for a row to be included in the output

(jan1 <- filter(flights, month==1 , day == 1))

## # A tibble: 842 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515             2       830             819
## 2  2013     1     1     533             529             4       850             830
## 3  2013     1     1     542             540             2       923             850
## 4  2013     1     1     544             545            -1      1004            1022
```

```
## 5 2013 1 1 554 600 -6 812 837
## 6 2013 1 1 554 558 -4 740 728
## 7 2013 1 1 555 600 -5 913 854
## 8 2013 1 1 557 600 -3 709 723
## 9 2013 1 1 557 600 -3 838 846
## 10 2013 1 1 558 600 -2 753 745
## # ... with 832 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>

# use near for double value equalities
near(sqrt(2) ^ 2, 2)
```

```
## [1] TRUE
```

```
(jan1 <- filter(flights, month==11 | month == 12))
```

```
## # A tibble: 55,403 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1 2013    11     1       5           2359           6       352           345
## 2 2013    11     1      35           2250          105       123           2356
## 3 2013    11     1     455           500           -5       641           651
## 4 2013    11     1     539           545           -6       856           827
## 5 2013    11     1     542           545           -3       831           855
## 6 2013    11     1     549           600          -11       912           923
## 7 2013    11     1     550           600          -10       705           659
## 8 2013    11     1     554           600           -6       659           701
## 9 2013    11     1     554           600           -6       826           827
## 10 2013    11     1     554           600           -6       749           751
## # ... with 55,393 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# equivalently
jan1 <- filter(flights, month %in% c(11,12))

# filter excludes both FALSE and NA values
df <- tibble(x = c(1, NA, 3))
filter(df, x>1)
```

```
## # A tibble: 1 x 1
##       x
##   <dbl>
## 1     3

filter(df, is.na(x) | x > 1)
```

```
## # A tibble: 2 x 1
##       x
##   <dbl>
## 1    NA
## 2     3
```

```
# Exercise 5.2.4
# Find all flights that:
# Had an arrival delay of two or more hours
filter(flights, arr_delay >= 2)
```

```
## # A tibble: 127,929 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1     517           515         2     830           819
## 2  2013     1     1     533           529         4     850           830
## 3  2013     1     1     542           540         2     923           850
## 4  2013     1     1     554           558        -4     740           728
## 5  2013     1     1     555           600        -5     913           854
## 6  2013     1     1     558           600        -2     753           745
## 7  2013     1     1     558           600        -2     924           917
## 8  2013     1     1     559           600        -1     941           910
## 9  2013     1     1     600           600         0     837           825
## 10 2013     1     1     602           605        -3     821           805
## # ... with 127,919 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
#Flew to Houston (IAH or HOU)
filter(flights, dest %in% c("IAH", "HOU"))
```

```
## # A tibble: 9,313 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1     517           515         2     830           819
## 2  2013     1     1     533           529         4     850           830
## 3  2013     1     1     623           627        -4     933           932
## 4  2013     1     1     728           732        -4    1041          1038
## 5  2013     1     1     739           739         0    1104          1038
## 6  2013     1     1     908           908         0    1228          1219
## 7  2013     1     1    1028          1026         2    1350          1339
## 8  2013     1     1    1044          1045        -1    1352          1351
## 9  2013     1     1    1114           900        134    1447          1222
## 10 2013     1     1    1205          1200         5    1503          1505
## # ... with 9,303 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
#Were operated by United, American, or Delta
filter(flights, carrier %in% c("UA", "AM", "DEL"))
```

```
## # A tibble: 58,665 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1     517           515         2     830           819
## 2  2013     1     1     533           529         4     850           830
## 3  2013     1     1     554           558        -4     740           728
## 4  2013     1     1     558           600        -2     924           917
## 5  2013     1     1     558           600        -2     923           937
## 6  2013     1     1     559           600        -1     854           902
## 7  2013     1     1     607           607         0     858           915
## 8  2013     1     1     611           600        11     945           931
## 9  2013     1     1     623           627        -4     933           932
## 10 2013     1     1     628           630        -2    1016          947
## # ... with 58,655 more rows, and 11 more variables: arr_delay <dbl>,
```

```
## # carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## # air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# Departed in summer (July, August, and September)
filter(flights, month %in% c(7, 8, 9))
```

```
## # A tibble: 86,326 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     7     1       1           2029        212      236           2359
## 2  2013     7     1       2           2359         3       344           344
## 3  2013     7     1      29           2245       104      151            1
## 4  2013     7     1      43           2130       193      322            14
## 5  2013     7     1      44           2150       174      300           100
## 6  2013     7     1      46           2051       235      304           2358
## 7  2013     7     1      48           2001       287      308           2305
## 8  2013     7     1      58           2155       183      335            43
## 9  2013     7     1     100           2146       194      327            30
## 10 2013     7     1     100           2245       135      337           135
```

```
## # ... with 86,316 more rows, and 11 more variables: arr_delay <dbl>,
## # carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## # air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# Arrived more than two hours late, but didn't leave late
filter(flights, dep_delay <= 0 & arr_delay >= 120)
```

```
## # A tibble: 29 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     1    27    1419           1420        -1     1754           1550
## 2  2013    10     7    1350           1350         0     1736           1526
## 3  2013    10     7    1357           1359        -2     1858           1654
## 4  2013    10    16     657           700        -3     1258           1056
## 5  2013    11     1     658           700        -2     1329           1015
## 6  2013     3    18    1844           1847        -3         39           2219
## 7  2013     4    17    1635           1640        -5     2049           1845
## 8  2013     4    18     558           600        -2     1149           850
## 9  2013     4    18     655           700        -5     1213           950
## 10 2013     5    22    1827           1830        -3     2217           2010
```

```
## # ... with 19 more rows, and 11 more variables: arr_delay <dbl>, carrier <chr>,
## # flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## # distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
flights[between(flights$dep_delay, 0, 120), ]
```

```
## # A tibble: 143,478 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     1     1     517           515         2     830           819
## 2  2013     1     1     533           529         4     850           830
## 3  2013     1     1     542           540         2     923           850
## 4  2013     1     1     559           559         0     702           706
## 5  2013     1     1     600           600         0     851           858
## 6  2013     1     1     600           600         0     837           825
## 7  2013     1     1     601           600         1     844           850
## 8  2013     1     1     607           607         0     858           915
```

```
## 9 2013 1 1 608 600 8 807 735
## 10 2013 1 1 611 600 11 945 931
## # ... with 143,468 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# Were delayed by at least an hour, but made up over 30 minutes in flight
filter(flights, dep_delay > 60 & arr_delay <= 30)
```

```
## # A tibble: 211 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1 2013     1     3    1850           1745         65     2148           2120
## 2 2013     1     3    1950           1845         65     2228           2227
## 3 2013     1     6    1019           900          79     1558           1530
## 4 2013     1     7    1543           1430         73     1758           1735
## 5 2013     1    12    1706           1600         66     1949           1927
## 6 2013     1    12    1953           1845         68     2154           2137
## 7 2013     1    19    1456           1355         61     1636           1615
## 8 2013     1    21    1531           1430         61     1843           1815
## 9 2013     1    21    1648           1545         63     1939           1910
## 10 2013    10     5    1605           1500         65     1857           1827
## # ... with 201 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# Departed between midnight and 6am (inclusive)
filter(flights, dep_time <= 600)
```

```
## # A tibble: 9,344 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1 2013     1     1     517           515          2      830           819
## 2 2013     1     1     533           529          4      850           830
## 3 2013     1     1     542           540          2      923           850
## 4 2013     1     1     544           545         -1     1004           1022
## 5 2013     1     1     554           600         -6      812           837
## 6 2013     1     1     554           558         -4      740           728
## 7 2013     1     1     555           600         -5      913           854
## 8 2013     1     1     557           600         -3      709           723
## 9 2013     1     1     557           600         -3      838           846
## 10 2013     1     1     558           600         -2      753           745
## # ... with 9,334 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
#How many flights have a missing dep_time
paste0(nrow(filter(flights, is.na(dep_time))), " flights have missing dep_time value")
```

```
## [1] "8255 flights have missing dep_time value"
```

```
# complete.cases gives TRUE when all values in a row are not NA
flights[!complete.cases(flights), ]
```

```
## # A tibble: 9,430 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
```

```
## 1 2013 1 1 1525 1530 -5 1934 1805
## 2 2013 1 1 1528 1459 29 2002 1647
## 3 2013 1 1 1740 1745 -5 2158 2020
## 4 2013 1 1 1807 1738 29 2251 2103
## 5 2013 1 1 1939 1840 59 29 2151
## 6 2013 1 1 1952 1930 22 2358 2207
## 7 2013 1 1 2016 1930 46 NA 2220
## 8 2013 1 1 NA 1630 NA NA 1815
## 9 2013 1 1 NA 1935 NA NA 2240
## 10 2013 1 1 NA 1500 NA NA 1825
## # ... with 9,420 more rows, and 11 more variables: arr_delay <dbl>,
## # carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## # air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

columns that have NA

```
colnames(flights[!complete.cases(flights), ])
```

```
## [1] "year" "month" "day" "dep_time"
## [5] "sched_dep_time" "dep_delay" "arr_time" "sched_arr_time"
## [9] "arr_delay" "carrier" "flight" "tailnum"
## [13] "origin" "dest" "air_time" "distance"
## [17] "hour" "minute" "time_hour"
```

```
?complete.cases
```

Arrange rows by getting a set of column names (or more complicated expressions) to order by
`arrange(flights, year, month, desc(day))`

```
## # A tibble: 336,776 x 19
##   year month day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int> <int> <int> <dbl> <int> <int>
## 1 2013 1 31 1 2100 181 124 2225
## 2 2013 1 31 4 2359 5 455 444
## 3 2013 1 31 7 2359 8 453 437
## 4 2013 1 31 12 2250 82 132 7
## 5 2013 1 31 26 2154 152 328 50
## 6 2013 1 31 34 2159 155 135 2315
## 7 2013 1 31 37 2249 108 132 2357
## 8 2013 1 31 54 2250 124 152 2359
## 9 2013 1 31 453 500 -7 651 648
## 10 2013 1 31 522 525 -3 820 820
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## # carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## # air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Missing values are always sorted at the end:

```
df <- tibble(x = c(5,2,NA))
arrange(df, x)
```

```
## # A tibble: 3 x 1
##   x
##   <dbl>
## 1 2
## 2 5
## 3 NA
```

5.3.1 Exercises

```
# Sort missing values in the start
arrange(df, desc(is.na(x)), x)
```

```
## # A tibble: 3 x 1
##       x
##   <dbl>
## 1    NA
## 2     2
## 3     5
```

```
# Sort flights to find the most delayed flights.
arrange(flights, desc(dep_delay))[1,]
```

```
## # A tibble: 1 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
## 1  2013     1     9     641             900     1301    1242         1530
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# Find the flights that left earliest.
arrange(flights, dep_time)[1,]
```

```
## # A tibble: 1 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
## 1  2013     1    13         1             2249      72     108         2357
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# Sort flights to find the fastest flights.
arrange(flights, distance/air_time)[1,]
```

```
## # A tibble: 1 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
## 1  2013     1    28    1917             1825      52    2118         1935
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# Which travelled the shortest?
arrange(flights, distance)[1,]
```

```
## # A tibble: 1 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
## 1  2013     7    27      NA             106      NA     NA         245
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# Which flights travelled the farthest?
```

```
arrange(flights, desc(distance))[1,]
```

```
## # A tibble: 1 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     1     1     857           900        -3     1516           1530
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
library(nycflights13)
```

```
library(tidyverse)
```

```
select(flights, year, month, day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

```
# Select all columns between year and day (inclusive)
```

```
select(flights, year:day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

```
# Select all columns except those from year to day (inclusive)
```

```
select(flights, -(year:day))
```

```
## # A tibble: 336,776 x 16
##   dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier
##   <int>         <int>      <dbl>    <int>         <int>      <dbl> <chr>
## 1     517           515         2      830           819        11 UA
```



```
## 2      533      529      4      850      830      20 UA
## 3      542      540      2      923      850      33 AA
## 4      544      545     -1     1004     1022     -18 B6
## 5      554      600     -6      812      837     -25 DL
## 6      554      558     -4      740      728      12 UA
## 7      555      600     -5      913      854      19 B6
## 8      557      600     -3      709      723     -14 EV
## 9      557      600     -3      838      846      -8 B6
## 10     558      600     -2      753      745       8 AA
## # ... with 336,766 more rows, and 9 more variables: flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# select all columns whose name contains string "ela"
select(flights, contains("ela"))
```

```
## # A tibble: 336,776 x 2
##   dep_delay arr_delay
##   <dbl>      <dbl>
## 1         2         11
## 2         4         20
## 3         2         33
## 4        -1        -18
## 5        -6        -25
## 6        -4         12
## 7        -5         19
## 8        -3        -14
## 9        -3         -8
## 10       -2          8
## # ... with 336,766 more rows
```

```
# select all columns whose name ends with "_time"
select(flights, ends_with("_time"))
```

```
## # A tibble: 336,776 x 5
##   dep_time sched_dep_time arr_time sched_arr_time air_time
##   <int>      <int>      <int>      <int>      <dbl>
## 1      517         515      830         819      227
## 2      533         529      850         830      227
## 3      542         540      923         850      160
## 4      544         545     1004        1022      183
## 5      554         600      812         837      116
## 6      554         558      740         728      150
## 7      555         600      913         854      158
## 8      557         600      709         723       53
## 9      557         600      838         846      140
## 10     558         600      753         745      138
## # ... with 336,766 more rows
```

```
# rename a column
rename(flights, departure_time=dep_time, arrival_time = arr_time)
```

```
## # A tibble: 336,776 x 19
##   year month day departure_time sched_dep_time dep_delay arrival_time
##   <int> <int> <int>      <int>      <int>      <dbl>      <int>
## 1  2013     1     1         517         515          2         830
```

```
## 2 2013 1 1 533 529 4 850
## 3 2013 1 1 542 540 2 923
## 4 2013 1 1 544 545 -1 1004
## 5 2013 1 1 554 600 -6 812
## 6 2013 1 1 554 558 -4 740
## 7 2013 1 1 555 600 -5 913
## 8 2013 1 1 557 600 -3 709
## 9 2013 1 1 557 600 -3 838
## 10 2013 1 1 558 600 -2 753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
## #   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

```
# handful of variables you'd like to move to the start of the data frame.
select(flights, dep_time, arr_time, sched_dep_time, sched_arr_time, everything())
```

```
## # A tibble: 336,776 x 19
##   dep_time arr_time sched_dep_time sched_arr_time year month day dep_delay
##   <int> <int> <int> <int> <int> <int> <int> <dbl>
## 1 517 830 515 819 2013 1 1 2
## 2 533 850 529 830 2013 1 1 4
## 3 542 923 540 850 2013 1 1 2
## 4 544 1004 545 1022 2013 1 1 -1
## 5 554 812 600 837 2013 1 1 -6
## 6 554 740 558 728 2013 1 1 -4
## 7 555 913 600 854 2013 1 1 -5
## 8 557 709 600 723 2013 1 1 -3
## 9 557 838 600 846 2013 1 1 -3
## 10 558 753 600 745 2013 1 1 -2
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

5.4.1 Exercises

```
# Brainstorm as many ways as possible to select dep_time, dep_delay, arr_time, and arr_delay from flights
```

```
select(flights, dep_time, dep_delay, arr_time, arr_delay)
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int> <dbl> <int> <dbl>
## 1 517 2 830 11
## 2 533 4 850 20
## 3 542 2 923 33
## 4 544 -1 1004 -18
## 5 554 -6 812 -25
## 6 554 -4 740 12
## 7 555 -5 913 19
## 8 557 -3 709 -14
## 9 557 -3 838 -8
## 10 558 -2 753 8
## # ... with 336,766 more rows
```

```
select(flights, starts_with("dep_") | starts_with("arr_"))
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>     <dbl>   <int>     <dbl>
## 1      517         2     830         11
## 2      533         4     850         20
## 3      542         2     923         33
## 4      544        -1    1004        -18
## 5      554        -6     812        -25
## 6      554        -4     740         12
## 7      555        -5     913         19
## 8      557        -3     709        -14
## 9      557        -3     838         -8
## 10     558        -2     753          8
## # ... with 336,766 more rows
```

```
select(flights, starts_with("dep_") & (ends_with("_delay") | ends_with("_time"))
      | starts_with("arr_") & (ends_with("_delay") | ends_with("_time")))
```

```
## # A tibble: 336,776 x 4
##   dep_delay dep_time arr_delay arr_time
##   <dbl>     <int>   <dbl>   <int>
## 1         2     517         11     830
## 2         4     533         20     850
## 3         2     542         33     923
## 4        -1     544        -18    1004
## 5        -6     554        -25     812
## 6        -4     554         12     740
## 7        -5     555         19     913
## 8        -3     557        -14     709
## 9        -3     557         -8     838
## 10       -2     558          8     753
## # ... with 336,766 more rows
```

What happens if you include the name of a variable multiple times in a select() call?

```
select(flights, dep_time, dep_time, dep_time, arr_delay)
```

```
## # A tibble: 336,776 x 2
##   dep_time arr_delay
##   <int>     <dbl>
## 1      517         11
## 2      533         20
## 3      542         33
## 4      544        -18
## 5      554        -25
## 6      554         12
## 7      555         19
## 8      557        -14
## 9      557         -8
## 10     558          8
## # ... with 336,766 more rows
```

What does the one_of() function do? Why might it be helpful in conjunction with this vector?

```
cols <- c("dep_time", "XXX", "dep_delay", "ZZZ", "QQQ", "arr_delay")
```

```
select(flights, one_of(cols))
```

```
## Warning: Unknown columns: `XXX`, `ZZZ`, `QQQ`
```

```
## # A tibble: 336,776 x 3
```

```
##   dep_time dep_delay arr_delay
```

```
##   <int>      <dbl>      <dbl>
```

```
## 1      517         2        11
```

```
## 2      533         4        20
```

```
## 3      542         2        33
```

```
## 4      544        -1       -18
```

```
## 5      554        -6       -25
```

```
## 6      554        -4        12
```

```
## 7      555        -5        19
```

```
## 8      557        -3       -14
```

```
## 9      557        -3        -8
```

```
## 10     558        -2         8
```

```
## # ... with 336,766 more rows
```

```
# case insensetivity is surprising
```

```
select(flights, contains("TIME", ignore.case = F))
```

```
## # A tibble: 336,776 x 0
```

```
library(nycflights13)
```

```
library(tidyverse)
```

```
# mutate() always adds new columns at the end of your dataset so we'll start by creating a narrower dat
```

```
flights_small <- select(flights, year:day, ends_with("delay"), distance, air_time)
```

```
# view(flights_small)
```

```
flights_small
```

```
## # A tibble: 336,776 x 7
```

```
##   year month   day dep_delay arr_delay distance air_time
```

```
##   <int> <int> <int>      <dbl>      <dbl>      <dbl>      <dbl>
```

```
## 1  2013     1     1         2        11      1400       227
```

```
## 2  2013     1     1         4        20      1416       227
```

```
## 3  2013     1     1         2        33      1089       160
```

```
## 4  2013     1     1        -1       -18      1576       183
```

```
## 5  2013     1     1        -6       -25       762       116
```

```
## 6  2013     1     1        -4        12       719       150
```

```
## 7  2013     1     1        -5        19      1065       158
```

```
## 8  2013     1     1        -3       -14       229        53
```

```
## 9  2013     1     1        -3        -8       944       140
```

```
## 10 2013     1     1        -2         8       733       138
```

```
## # ... with 336,766 more rows
```

```
mutate(flights, gain = dep_delay - arr_delay, gain_per_hour = gain/hour,  
       speed = distance/air_time *60)
```

```
## # A tibble: 336,776 x 22
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
```

```
##   <int> <int> <int>   <int>      <int>      <dbl>      <int>      <int>
```

```
## 1  2013     1     1     517         515         2      830        819
```

```
## 2 2013 1 1 533 529 4 850 830
## 3 2013 1 1 542 540 2 923 850
## 4 2013 1 1 544 545 -1 1004 1022
## 5 2013 1 1 554 600 -6 812 837
## 6 2013 1 1 554 558 -4 740 728
## 7 2013 1 1 555 600 -5 913 854
## 8 2013 1 1 557 600 -3 709 723
## 9 2013 1 1 557 600 -3 838 846
## 10 2013 1 1 558 600 -2 753 745
## # ... with 336,766 more rows, and 14 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>,
## #   gain <dbl>, gain_per_hour <dbl>, speed <dbl>
```

If you only want to keep the new variables, use transmute()

```
transmute(flights, gain = dep_delay - arr_delay, gain_per_hour = gain/hour,
          speed = distance/air_time *60)
```

```
## # A tibble: 336,776 x 3
##   gain gain_per_hour speed
##   <dbl>      <dbl> <dbl>
## 1    -9        -1.8  370.
## 2   -16        -3.2  374.
## 3   -31        -6.2  408.
## 4    17         3.4  517.
## 5    19         3.17 394.
## 6   -16        -3.2  288.
## 7   -24         -4   404.
## 8    11         1.83  259.
## 9     5         0.833 405.
## 10  -10        -1.67  319.
## # ... with 336,766 more rows
```

function must be vectorised to be able to use in mutate

Modular arithmetic: %/% (integer division) and %% (remainder) are vectorized

```
transmute(flights, hour = dep_time %/% 100, min = dep_time %% 100)
```

```
## # A tibble: 336,776 x 2
##   hour min
##   <dbl> <dbl>
## 1     5  17
## 2     5  33
## 3     5  42
## 4     5  44
## 5     5  54
## 6     5  54
## 7     5  55
## 8     5  57
## 9     5  57
## 10    5  58
## # ... with 336,766 more rows
```

log() functions are also vectorize

```
transmute(flights, gain = dep_delay - arr_delay, gain_per_hour = gain/hour, logOfGain = log2(gain_per_h
```

```
## Warning: NaNs produced
## # A tibble: 336,776 x 3
##   gain gain_per_hour logOfGain
##   <dbl>         <dbl>     <dbl>
## 1    -9          -1.8      NaN
## 2   -16          -3.2      NaN
## 3   -31          -6.2      NaN
## 4    17           3.4       1.77
## 5    19           3.17      1.66
## 6   -16          -3.2      NaN
## 7   -24           -4       NaN
## 8    11           1.83      0.874
## 9     5           0.833     -0.263
## 10  -10          -1.67      NaN
## # ... with 336,766 more rows

# lead() and lag() allows you to compute running differences (e.g. x - lag(x)) or
# find when values change (x != lag(x)).

(x <- 1:10)

## [1] 1 2 3 4 5 6 7 8 9 10
lag(x)

## [1] NA 1 2 3 4 5 6 7 8 9
lead(x)

## [1] 2 3 4 5 6 7 8 9 10 NA
transmute(flights, dest, lag(dest), lead(dest))

## # A tibble: 336,776 x 3
##   dest `lag(dest)` `lead(dest)`
##   <chr> <chr>      <chr>
## 1 IAH   <NA>        IAH
## 2 IAH   IAH         MIA
## 3 MIA   IAH         BQN
## 4 BQN   MIA         ATL
## 5 ATL   BQN         ORD
## 6 ORD   ATL         FLL
## 7 FLL   ORD         IAD
## 8 IAD   FLL         MCO
## 9 MCO   IAD         ORD
## 10 ORD  MCO         PBI
## # ... with 336,766 more rows

# Cumulative and rolling aggregates (i.e. a sum computed over a rolling window):
# R => cumsum(), cumprod(), cummin(), cummax();
(x <- 1:10)

## [1] 1 2 3 4 5 6 7 8 9 10
cumsum(x)

## [1] 1 3 6 10 15 21 28 36 45 55
```

```
cumprod(x)
```

```
## [1]      1      2      6     24    120    720   5040  40320 362880
## [10] 3628800
```

```
cummin(x)
```

```
## [1] 1 1 1 1 1 1 1 1 1 1
```

```
cummax(x)
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
# dplyr => cummean() for cumulative means
```

```
cummean(x)
```

```
## [1] 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0 5.5
```

```
# For Rolling aggregates use RcppRoll package
```

```
x <- matrix(rnorm(100),nrow=50,ncol=2)
```

```
x
```

```
##           [,1]      [,2]
## [1,] -0.29959524 -0.63147089
## [2,] -0.32246643  1.38927160
## [3,]  0.39175228 -0.70178125
## [4,] -0.65047246 -0.29528673
## [5,] -0.31543296  0.01803912
## [6,] -1.74899113  1.20615087
## [7,] -0.90550442 -0.18016485
## [8,] -1.63085485 -0.79654724
## [9,] -0.01186019  0.82201281
## [10,]  0.51400268 -0.43872827
## [11,]  0.32506658 -0.77086793
## [12,] -0.52904793  0.89440511
## [13,] -2.08637920 -0.18307730
## [14,] -0.49910263  0.57321819
## [15,] -0.71880337  0.58602363
## [16,]  0.33800372  0.89548646
## [17,] -1.87720006 -1.82684575
## [18,]  1.01635860  1.58421898
## [19,]  0.25625602  0.98569845
## [20,] -0.42996769  1.04024702
## [21,]  2.37215561  0.61679673
## [22,] -0.77185834  0.21951974
## [23,]  1.95081160  1.29609260
## [24,] -1.41486967 -0.06892722
## [25,]  0.20850842  0.77680706
## [26,]  0.70111753 -1.29330171
## [27,] -0.59051599  1.17481428
## [28,] -1.17618094 -0.25125114
## [29,] -0.83247354 -0.76162277
## [30,] -1.20257030  0.79582482
## [31,] -1.97827664 -0.33878024
## [32,] -0.56421367 -1.06955657
## [33,]  1.17510767 -0.19830257
## [34,]  0.75657449 -0.62932429
```

```
## [35,] 0.88684945 -1.06833687
## [36,] -0.18783132 0.84877298
## [37,] 0.61024189 -0.56139558
## [38,] -0.36923513 0.85426007
## [39,] -0.68913266 1.22024214
## [40,] -0.15747136 0.72666180
## [41,] -0.97592204 -0.20249386
## [42,] 1.46461134 0.55606345
## [43,] -0.77340892 0.98374893
## [44,] 0.64020765 0.19751422
## [45,] 0.81140438 -1.08392045
## [46,] 0.07058968 1.48854290
## [47,] 0.67038346 0.37198268
## [48,] -1.14539236 -1.81877756
## [49,] -2.85036021 -1.20865887
## [50,] 0.03711969 0.77637460
```

```
RcppRoll::roll_sum(x,12)
```

```
##           [,1]      [,2]
## [1,] -5.18340408 0.5150323
## [2,] -6.97018803 0.9634259
## [3,] -7.14682423 0.1473725
## [4,] -8.25737988 1.4351774
## [5,] -7.26890370 2.6259506
## [6,] -8.83067080 0.7810657
## [7,] -6.06532107 1.1591338
## [8,] -4.90356063 2.3249971
## [9,] -3.70267347 4.1617914
## [10,] -1.31865766 3.9565753
## [11,] -2.60451869 4.6148233
## [12,] -0.97877366 6.6817839
## [13,] -1.86459540 5.7184515
## [14,] 0.43029221 6.6783359
## [15,] 1.63051237 4.8118160
## [16,] 1.75879975 5.4006066
## [17,] 0.24461508 4.2538690
## [18,] 1.28934161 5.3190920
## [19,] -0.92958729 4.5306978
## [20,] -3.16411995 3.2062192
## [21,] -3.29836593 1.0964156
## [22,] -4.49541387 0.2813163
## [23,] -2.96698104 -0.5675278
## [24,] -4.03094319 -2.9319572
## [25,] -2.80390483 -2.0142570
## [26,] -2.40217136 -3.3524597
## [27,] -3.47252402 -1.2048979
## [28,] -3.57114069 -1.1594700
## [29,] -2.55243111 -0.1815571
## [30,] -2.69587962 0.3775718
## [31,] -0.02869798 0.1378104
## [32,] 1.17616974 1.4603396
## [33,] 2.38059106 2.7274104
## [34,] 2.01688777 1.8417925
## [35,] 1.33090296 3.9596597
```



```
## [36,] 1.11443697 5.3999793
## [37,] 0.15687593 2.7324287
## [38,] -3.30372617 2.0851655
## [39,] -2.89737135 2.0072800

# ranking (basically report the indices of elements if they were sorted, NA and INF goes to the end)
x <- c(5, 1, 3, Inf, 2, 2, NA) # => (1,2,2,3,5,Inf,NA)
row_number(x)

## [1] 5 1 4 6 2 3 NA
min_rank(x)

## [1] 5 1 4 6 2 2 NA
dense_rank(x)

## [1] 4 1 3 5 2 2 NA
percent_rank(x) # a number between 0 and 1 computed by rescaling min_rank to [0, 1]

## [1] 0.8 0.0 0.6 1.0 0.2 0.2 NA
cume_dist(x) # a cumulative distribution function. Proportion of all values less than or equal to the c

## [1] 0.8333333 0.1666667 0.6666667 1.0000000 0.5000000 0.5000000 NA
# ntile creates a rough rank, which breaks the input vector into n buckets
ntile(x, 2)

## [1] 2 1 2 2 1 1 NA
ntile(runif(100), 10)

## [1] 5 9 2 1 3 10 10 10 4 2 9 7 5 6 5 9 4 1 6 10 8 9 2 2 1
## [26] 4 10 8 3 10 4 5 10 5 4 2 5 8 3 2 10 8 6 3 3 2 1 2 10 5
## [51] 6 6 6 10 9 8 2 9 8 3 9 5 9 9 7 9 7 3 2 7 8 5 7 1 3
## [76] 6 4 7 5 1 4 8 7 6 7 4 4 8 3 1 7 3 7 8 4 6 1 1 6 1

# 5.5.2 Exercises

# Convert dep_time and sched_dep_time to a more convenient representation of number of minutes since mi

# flights[order(flights$dep_time, na.last = T, decreasing = T), ]
flights1 <- mutate(flights, dep_time_minute = (dep_time%/%100) * 60 + dep_time%%100, sched_dep_time_min
select (flights1, year, month, day, dep_time, dep_time_minute, sched_dep_time, sched_dep_time_minute, every

## # A tibble: 336,776 x 21
##   year month   day dep_time dep_time_minute sched_dep_time sched_dep_time_~
##   <int> <int> <int>   <int>         <dbl>         <int>         <dbl>
## 1  2013     1     1     517           317           515           315
## 2  2013     1     1     533           333           529           329
## 3  2013     1     1     542           342           540           340
## 4  2013     1     1     544           344           545           345
## 5  2013     1     1     554           354           600           360
## 6  2013     1     1     554           354           558           358
## 7  2013     1     1     555           355           600           360
## 8  2013     1     1     557           357           600           360
## 9  2013     1     1     557           357           600           360
## 10 2013     1     1     558           358           600           360
```

```
## # ... with 336,766 more rows, and 14 more variables: dep_delay <dbl>,
## #   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## #   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# Compare air_time with arr_time - dep_time. What do you expect to see? What do you see? What do you ne
fixed_air_time_flights <- transmute(flights, arr_time, dep_time, air_time, fixed_air_time_minute = abs
fixed_air_time_flights[order(fixed_air_time_flights$fixed_air_time), ]
```

```
## # A tibble: 336,776 x 5
##   arr_time dep_time air_time fixed_air_time_minute fixed_air_time
##   <int>    <int>    <dbl>          <dbl>          <dbl>
## 1    1206    1133     23             33             33
## 2    1358    1323     23             35             35
## 3    1347    1312     23             35             35
## 4    1238    1203     21             35             35
## 5    1531    1455     22             36             36
## 6     758     722     22             36             36
## 7     758     722     22             36             36
## 8     754     718     24             36             36
## 9    1403    1326     22             37             37
## 10   1533    1456     21             37             37
## # ... with 336,766 more rows
```

```
filter(fixed_air_time_flights, arr_time <= dep_time)
```

```
## # A tibble: 10,633 x 5
##   arr_time dep_time air_time fixed_air_time_minute fixed_air_time
##   <int>    <int>    <dbl>          <dbl>          <dbl>
## 1      3    1929     192             1166           1926
## 2     29    1939      NA             1150           1910
## 3      8    2058     159             1250           2050
## 4    146    2102     199             1156           1916
## 5     25    2108     354             1243           2043
## 6     16    2120     160             1264           2104
## 7      6    2121     143             1275           2115
## 8     26    2128     338             1262           2102
## 9     20    2134     152             1274           2114
## 10    25    2136     154             1271           2111
## # ... with 10,623 more rows
```

```
# Compare dep_time, sched_dep_time, and dep_delay. How would you expect those three numbers to be relat
transmute(flights, dep_time, sched_dep_time, dep_delay, dep_delay_fixed = dep_time - sched_dep_time )
```

```
## # A tibble: 336,776 x 4
##   dep_time sched_dep_time dep_delay dep_delay_fixed
##   <int>         <int>    <dbl>         <int>
## 1     517           515         2             2
## 2     533           529         4             4
## 3     542           540         2             2
## 4     544           545        -1            -1
## 5     554           600        -6           -46
## 6     554           558        -4            -4
## 7     555           600        -5           -45
## 8     557           600        -3           -43
## 9     557           600        -3           -43
```

```
## 10      558      600      -2      -42
## # ... with 336,766 more rows
```

```
# Find the 10 most delayed flights using a ranking function. How do you want to handle ties? Carefully
filter(flights, min_rank(desc(dep_delay)) <= 10)
```

```
## # A tibble: 10 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     9     641             900         1301    1242         1530
## 2  2013     1    10    1121            1635         1126    1239         1810
## 3  2013    12     5     756            1700          896    1058         2020
## 4  2013     3    17    2321             810          911     135         1020
## 5  2013     4    10    1100            1900          960    1342         2211
## 6  2013     6    15    1432            1935         1137    1607         2120
## 7  2013     6    27     959            1900          899    1236         2226
## 8  2013     7    22     845            1600         1005    1044         1815
## 9  2013     7    22    2257             759          898     121         1026
## 10 2013     9    20    1139            1845         1014    1457         2210
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
flights[order(flights$dep_delay, decreasing = T), ]
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     9     641             900         1301    1242         1530
## 2  2013     6    15    1432            1935         1137    1607         2120
## 3  2013     1    10    1121            1635         1126    1239         1810
## 4  2013     9    20    1139            1845         1014    1457         2210
## 5  2013     7    22     845            1600         1005    1044         1815
## 6  2013     4    10    1100            1900          960    1342         2211
## 7  2013     3    17    2321             810          911     135         1020
## 8  2013     6    27     959            1900          899    1236         2226
## 9  2013     7    22    2257             759          898     121         1026
## 10 2013    12     5     756            1700          896    1058         2020
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
library(nycflights13)
library(tidyverse)
```

```
# summarize() with group_by() changes the unit of analysis from the complete dataset to individual groups
# Then, when you use the dplyr verbs on a grouped data frame they'll be automatically applied "by group"
```

```
(by_day <- group_by(flights, year, month, day))
```

```
## # A tibble: 336,776 x 19
## # Groups:   year, month, day [365]
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515          2     830         819
## 2  2013     1     1     533             529          4     850         830
```

```
## 3 2013 1 1 542 540 2 923 850
## 4 2013 1 1 544 545 -1 1004 1022
## 5 2013 1 1 554 600 -6 812 837
## 6 2013 1 1 554 558 -4 740 728
## 7 2013 1 1 555 600 -5 913 854
## 8 2013 1 1 557 600 -3 709 723
## 9 2013 1 1 557 600 -3 838 846
## 10 2013 1 1 558 600 -2 753 745
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## # carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## # air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
str(attributes(by_day))
```

```
## List of 4
## $ names : chr [1:19] "year" "month" "day" "dep_time" ...
## $ row.names: int [1:336776] 1 2 3 4 5 6 7 8 9 10 ...
## $ class : chr [1:4] "grouped_df" "tbl_df" "tbl" "data.frame"
## $ groups : tibble [365 x 4] (S3: tbl_df/tbl/data.frame)
## ..$ year : int [1:365] 2013 2013 2013 2013 2013 2013 2013 2013 2013 2013 ...
## ..$ month: int [1:365] 1 1 1 1 1 1 1 1 1 1 ...
## ..$ day : int [1:365] 1 2 3 4 5 6 7 8 9 10 ...
## ..$ .rows:List of 365
## .. ..$ : int [1:842] 1 2 3 4 5 6 7 8 9 10 ...
## .. ..$ : int [1:943] 843 844 845 846 847 848 849 850 851 852 ...
## .. ..$ : int [1:914] 1786 1787 1788 1789 1790 1791 1792 1793 1794 1795 ...
## .. ..$ : int [1:915] 2700 2701 2702 2703 2704 2705 2706 2707 2708 2709 ...
## .. ..$ : int [1:720] 3615 3616 3617 3618 3619 3620 3621 3622 3623 3624 ...
## .. ..$ : int [1:832] 4335 4336 4337 4338 4339 4340 4341 4342 4343 4344 ...
## .. ..$ : int [1:933] 5167 5168 5169 5170 5171 5172 5173 5174 5175 5176 ...
## .. ..$ : int [1:899] 6100 6101 6102 6103 6104 6105 6106 6107 6108 6109 ...
## .. ..$ : int [1:902] 6999 7000 7001 7002 7003 7004 7005 7006 7007 7008 ...
## .. ..$ : int [1:932] 7901 7902 7903 7904 7905 7906 7907 7908 7909 7910 ...
## .. ..$ : int [1:930] 8833 8834 8835 8836 8837 8838 8839 8840 8841 8842 ...
## .. ..$ : int [1:690] 9763 9764 9765 9766 9767 9768 9769 9770 9771 9772 ...
## .. ..$ : int [1:828] 10453 10454 10455 10456 10457 10458 10459 10460 10461 10462 ...
## .. ..$ : int [1:928] 11281 11282 11283 11284 11285 11286 11287 11288 11289 11290 ...
## .. ..$ : int [1:894] 12209 12210 12211 12212 12213 12214 12215 12216 12217 12218 ...
## .. ..$ : int [1:901] 13103 13104 13105 13106 13107 13108 13109 13110 13111 13112 ...
## .. ..$ : int [1:927] 14004 14005 14006 14007 14008 14009 14010 14011 14012 14013 ...
## .. ..$ : int [1:924] 14931 14932 14933 14934 14935 14936 14937 14938 14939 14940 ...
## .. ..$ : int [1:674] 15855 15856 15857 15858 15859 15860 15861 15862 15863 15864 ...
## .. ..$ : int [1:786] 16529 16530 16531 16532 16533 16534 16535 16536 16537 16538 ...
## .. ..$ : int [1:912] 17315 17316 17317 17318 17319 17320 17321 17322 17323 17324 ...
## .. ..$ : int [1:890] 18227 18228 18229 18230 18231 18232 18233 18234 18235 18236 ...
## .. ..$ : int [1:897] 19117 19118 19119 19120 19121 19122 19123 19124 19125 19126 ...
## .. ..$ : int [1:925] 20014 20015 20016 20017 20018 20019 20020 20021 20022 20023 ...
## .. ..$ : int [1:922] 20939 20940 20941 20942 20943 20944 20945 20946 20947 20948 ...
## .. ..$ : int [1:680] 21861 21862 21863 21864 21865 21866 21867 21868 21869 21870 ...
## .. ..$ : int [1:823] 22541 22542 22543 22544 22545 22546 22547 22548 22549 22550 ...
## .. ..$ : int [1:923] 23364 23365 23366 23367 23368 23369 23370 23371 23372 23373 ...
## .. ..$ : int [1:890] 24287 24288 24289 24290 24291 24292 24293 24294 24295 24296 ...
## .. ..$ : int [1:900] 25177 25178 25179 25180 25181 25182 25183 25184 25185 25186 ...
## .. ..$ : int [1:928] 26077 26078 26079 26080 26081 26082 26083 26084 26085 26086 ...
## .. ..$ : int [1:926] 111297 111298 111299 111300 111301 111302 111303 111304 111305 111306 ...
```

```

## .. ..$ : int [1:682] 112223 112224 112225 112226 112227 112228 112229 112230 112231 112232 ...
## .. ..$ : int [1:814] 112905 112906 112907 112908 112909 112910 112911 112912 112913 112914 ...
## .. ..$ : int [1:932] 113719 113720 113721 113722 113723 113724 113725 113726 113727 113728 ...
## .. ..$ : int [1:896] 114651 114652 114653 114654 114655 114656 114657 114658 114659 114660 ...
## .. ..$ : int [1:901] 115547 115548 115549 115550 115551 115552 115553 115554 115555 115556 ...
## .. ..$ : int [1:932] 116448 116449 116450 116451 116452 116453 116454 116455 116456 116457 ...
## .. ..$ : int [1:930] 117380 117381 117382 117383 117384 117385 117386 117387 117388 117389 ...
## .. ..$ : int [1:684] 118310 118311 118312 118313 118314 118315 118316 118317 118318 118319 ...
## .. ..$ : int [1:829] 118994 118995 118996 118997 118998 118999 119000 119001 119002 119003 ...
## .. ..$ : int [1:929] 119823 119824 119825 119826 119827 119828 119829 119830 119831 119832 ...
## .. ..$ : int [1:893] 120752 120753 120754 120755 120756 120757 120758 120759 120760 120761 ...
## .. ..$ : int [1:918] 121645 121646 121647 121648 121649 121650 121651 121652 121653 121654 ...
## .. ..$ : int [1:956] 122563 122564 122565 122566 122567 122568 122569 122570 122571 122572 ...
## .. ..$ : int [1:954] 123519 123520 123521 123522 123523 123524 123525 123526 123527 123528 ...
## .. ..$ : int [1:738] 124473 124474 124475 124476 124477 124478 124479 124480 124481 124482 ...
## .. ..$ : int [1:848] 125211 125212 125213 125214 125215 125216 125217 125218 125219 125220 ...
## .. ..$ : int [1:948] 126059 126060 126061 126062 126063 126064 126065 126066 126067 126068 ...
## .. ..$ : int [1:943] 127007 127008 127009 127010 127011 127012 127013 127014 127015 127016 ...
## .. ..$ : int [1:949] 127950 127951 127952 127953 127954 127955 127956 127957 127958 127959 ...
## .. ..$ : int [1:961] 128899 128900 128901 128902 128903 128904 128905 128906 128907 128908 ...
## .. ..$ : int [1:957] 129860 129861 129862 129863 129864 129865 129866 129867 129868 129869 ...
## .. ..$ : int [1:743] 130817 130818 130819 130820 130821 130822 130823 130824 130825 130826 ...
## .. ..$ : int [1:880] 131560 131561 131562 131563 131564 131565 131566 131567 131568 131569 ...
## .. ..$ : int [1:961] 132440 132441 132442 132443 132444 132445 132446 132447 132448 132449 ...
## .. ..$ : int [1:938] 133401 133402 133403 133404 133405 133406 133407 133408 133409 133410 ...
## .. ..$ : int [1:945] 134339 134340 134341 134342 134343 134344 134345 134346 134347 134348 ...
## .. ..$ : int [1:964] 135284 135285 135286 135287 135288 135289 135290 135291 135292 135293 ...
## .. ..$ : int [1:958] 136248 136249 136250 136251 136252 136253 136254 136255 136256 136257 ...
## .. ..$ : int [1:765] 137206 137207 137208 137209 137210 137211 137212 137213 137214 137215 ...
## .. ..$ : int [1:913] 137971 137972 137973 137974 137975 137976 137977 137978 137979 137980 ...
## .. ..$ : int [1:977] 138884 138885 138886 138887 138888 138889 138890 138891 138892 138893 ...
## .. ..$ : int [1:965] 139861 139862 139863 139864 139865 139866 139867 139868 139869 139870 ...
## .. ..$ : int [1:972] 140826 140827 140828 140829 140830 140831 140832 140833 140834 140835 ...
## .. ..$ : int [1:980] 141798 141799 141800 141801 141802 141803 141804 141805 141806 141807 ...
## .. ..$ : int [1:979] 142778 142779 142780 142781 142782 142783 142784 142785 142786 142787 ...
## .. ..$ : int [1:765] 143757 143758 143759 143760 143761 143762 143763 143764 143765 143766 ...
## .. ..$ : int [1:908] 144522 144523 144524 144525 144526 144527 144528 144529 144530 144531 ...
## .. ..$ : int [1:980] 145430 145431 145432 145433 145434 145435 145436 145437 145438 145439 ...
## .. ..$ : int [1:966] 146410 146411 146412 146413 146414 146415 146416 146417 146418 146419 ...
## .. ..$ : int [1:974] 147376 147377 147378 147379 147380 147381 147382 147383 147384 147385 ...
## .. ..$ : int [1:982] 148350 148351 148352 148353 148354 148355 148356 148357 148358 148359 ...
## .. ..$ : int [1:979] 149332 149333 149334 149335 149336 149337 149338 149339 149340 149341 ...
## .. ..$ : int [1:767] 150311 150312 150313 150314 150315 150316 150317 150318 150319 150320 ...
## .. ..$ : int [1:907] 151078 151079 151080 151081 151082 151083 151084 151085 151086 151087 ...
## .. ..$ : int [1:981] 151985 151986 151987 151988 151989 151990 151991 151992 151993 151994 ...
## .. ..$ : int [1:967] 152966 152967 152968 152969 152970 152971 152972 152973 152974 152975 ...
## .. ..$ : int [1:970] 153933 153934 153935 153936 153937 153938 153939 153940 153941 153942 ...
## .. ..$ : int [1:980] 154903 154904 154905 154906 154907 154908 154909 154910 154911 154912 ...
## .. ..$ : int [1:977] 155883 155884 155885 155886 155887 155888 155889 155890 155891 155892 ...
## .. ..$ : int [1:767] 156860 156861 156862 156863 156864 156865 156866 156867 156868 156869 ...
## .. ..$ : int [1:905] 157627 157628 157629 157630 157631 157632 157633 157634 157635 157636 ...
## .. ..$ : int [1:978] 158532 158533 158534 158535 158536 158537 158538 158539 158540 158541 ...
## .. ..$ : int [1:973] 159510 159511 159512 159513 159514 159515 159516 159517 159518 159519 ...
## .. ..$ : int [1:977] 160483 160484 160485 160486 160487 160488 160489 160490 160491 160492 ...

```

```
## .. ..$ : int [1:982] 161460 161461 161462 161463 161464 161465 161466 161467 161468 161469 ...
## .. ..$ : int [1:974] 162442 162443 162444 162445 162446 162447 162448 162449 162450 162451 ...
## .. ..$ : int [1:769] 163416 163417 163418 163419 163420 163421 163422 163423 163424 163425 ...
## .. ..$ : int [1:897] 164185 164186 164187 164188 164189 164190 164191 164192 164193 164194 ...
## .. ..$ : int [1:970] 165082 165083 165084 165085 165086 165087 165088 165089 165090 165091 ...
## .. ..$ : int [1:983] 166052 166053 166054 166055 166056 166057 166058 166059 166060 166061 ...
## .. ..$ : int [1:992] 167035 167036 167037 167038 167039 167040 167041 167042 167043 167044 ...
## .. ..$ : int [1:985] 168027 168028 168029 168030 168031 168032 168033 168034 168035 168036 ...
## .. ..$ : int [1:981] 169012 169013 169014 169015 169016 169017 169018 169019 169020 169021 ...
## .. ..$ : int [1:770] 169993 169994 169995 169996 169997 169998 169999 170000 170001 170002 ...
## .. ..$ : int [1:911] 170763 170764 170765 170766 170767 170768 170769 170770 170771 170772 ...
## .. ..$ : int [1:981] 171674 171675 171676 171677 171678 171679 171680 171681 171682 171683 ...
## .. ..$ : int [1:975] 172655 172656 172657 172658 172659 172660 172661 172662 172663 172664 ...
## .. .. [list output truncated]
## ..- attr(*, ".drop")= logi TRUE
```

```
# delay by day
summarise(by_day, delay = mean(dep_delay, na.rm = T))
```


```
## # A tibble: 365 x 4
## # Groups:   year, month [12]
##   year month   day delay
##   <int> <int> <int> <dbl>
## 1  2013     1     1  11.5
## 2  2013     1     2  13.9
## 3  2013     1     3  11.0
## 4  2013     1     4   8.95
## 5  2013     1     5   5.73
## 6  2013     1     6   7.15
## 7  2013     1     7   5.42
## 8  2013     1     8   2.55
## 9  2013     1     9   2.28
## 10 2013     1    10   2.84
## # ... with 355 more rows
```

```
library(nycflights13)
library(tidyverse)

# Explore the relationship between the distance and average delay for each location
delay <- flights %>%
  group_by(dest) %>%
  summarise(count = n(), dist = mean(distance, na.rm = T), delay = mean(arr_delay, na.rm = T)) %>%
  filter(count > 20 & dest != "HNL")

ggplot(data = delay, mapping = aes(x = dist, y = delay))+
  geom_point(aes(size=count), alpha=1/3)+
  geom_smooth(se=F)
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
# compare it with:
by_destination <- group_by(flights, dest)
# (avgDelayByDest <- summarise(by_destination, count = n(), dist = mean(distance, na.rm = T), delay = m
# (delay <- filter(avgDelayByDest, count > 20 & dest != "HNL"))
```

```
library(nycflights13)
library(tidyverse)

# count missing values in arr_delay
missing_arr_delay <- flights %>%
  filter(is.na(arr_delay) | is.na(dep_delay)) %>%
  group_by(arr_delay) %>%
  summarise(count = n())

sprintf("Number of records with missing arr_delay or dep_delay: %d",missing_arr_delay$count)
```

```
## [1] "Number of records with missing arr_delay or dep_delay: 9430"
```

```
# Remove records with missing arr_delay or dep_delay
(not_cancelled <- flights %>%
  filter(!is.na(arr_delay) & !is.na(dep_delay))
)
```

```
## # A tibble: 327,346 x 19
```


```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     1     1     517           515         2      830           819
## 2  2013     1     1     533           529         4      850           830
## 3  2013     1     1     542           540         2      923           850
```

```
## 4 2013 1 1 544 545 -1 1004 1022
## 5 2013 1 1 554 600 -6 812 837
## 6 2013 1 1 554 558 -4 740 728
## 7 2013 1 1 555 600 -5 913 854
## 8 2013 1 1 557 600 -3 709 723
## 9 2013 1 1 557 600 -3 838 846
## 10 2013 1 1 558 600 -2 753 745
## # ... with 327,336 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

let's look at the planes (identified by their tail number) that have the highest average delays:

```
delays <- not_cancelled %>%
  group_by(tailnum) %>%
  summarise(delay = mean(arr_delay))

ggplot(data = delays, mapping = aes(x = delay)) +
  geom_freqpoly(binwidth = 10)
```



Wow, there are some planes that have an average delay of 5 hours (300 minutes)!
Let's draw a scatterplot of number of flights vs. average delay:

```
(delays1 <- not_cancelled %>%
  group_by(tailnum) %>%
  summarise(count = n(), delay = mean(arr_delay))
)
```



```
## # A tibble: 4,037 x 3
##   tailnum count  delay
##   <chr>   <int> <dbl>
## 1 D942DN     4  31.5
## 2 NOEGMQ   352   9.98
## 3 N10156   145  12.7
## 4 N102UW    48   2.94
## 5 N103US    46  -6.93
## 6 N104UW    46   1.80
## 7 N10575   269  20.7
## 8 N105UW    45  -0.267
## 9 N107US    41  -5.73
## 10 N108UW    60  -1.25
## # ... with 4,027 more rows
```

```
# filter(delays1 , is.na(delay) | is.na(count))
```


```
delays1 %>%
```

```
# it's often useful to filter out the groups with the smallest numbers of observations to see pattern.
```

```
filter(count > 50) %>%
```

```
ggplot(mapping = aes(x=count, y=delay))+
```

```
geom_point(alpha=1/10)
```



```
# ANOTHER EXAMPLE:
```

```
# When I plot the skill of the batter (measured by the batting average, ba) against the number of oppor  
# (measured by at bat, ab), you see two patterns:
```

```
batting <- as_tibble(Lahman::Batting)
```

```
(batters <- batting %>%
```

```

group_by(playerID) %>%
  summarise(ba = sum(H, na.rm = T)/sum(AB, na.rm = T), ab = sum(AB, na.rm = T))
)


## # A tibble: 19,428 x 3
##   playerID      ba    ab
##   <chr>      <dbl> <int>
## 1 aardsda01  0         4
## 2 aaronha01  0.305    12364
## 3 aaronto01  0.229     944
## 4 aasedo01   0         5
## 5 abadan01   0.0952    21
## 6 abadfe01   0.111     9
## 7 abadijo01  0.224     49
## 8 abbated01  0.254    3044
## 9 abbeybe01  0.169     225
## 10 abbeych01 0.281    1756
## # ... with 19,418 more rows

# Variation in our aggregate decreases as we get more data points.
# There's a positive correlation between skill (ba) and opportunities to hit the ball (ab).

batters %>%
  filter (ab > 100) %>%
  ggplot(mapping = aes(x=ab, y = ba)) +
  geom_point()+
  geom_smooth(se=F)

## `geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'

```



If you naively sort on desc(ba), the people with the best batting averages are clearly lucky, not ski

```
batters %>%
  arrange(desc(ba))
```

```
## # A tibble: 19,428 x 3
##   playerID    ba    ab
##   <chr>    <dbl> <int>
## 1 abramge01     1     1
## 2 alberan01     1     1
## 3 allarko01     1     1
## 4 banisje01     1     1
## 5 bartocl01     1     1
## 6 bassdo01     1     1
## 7 birasst01     1     2
## 8 bruneju01     1     1
## 9 burnscb01     1     1
## 10 cammaer01     1     1
## # ... with 19,418 more rows
```

```
# sum(x > 10) can take a logical expression that filters in certain records and then it adds them (sum
# mean(x > 60) can filters in records whose column 'x' value is greater than 60 and then calculate th
# Measures of location: mean(x), median(x) (Half of the values of x is less than median(x) and other ha
# Measures of spread: sd(x), IQR(x), mad(x)
# Measures of rank: min(x), quantile(x, 0.25), max(x)
# Measures of position: first(x), nth(x, 2), last(x). These work similarly to x[1], x[2], and x[length(
# Counts:
```

```

# n() : returns the size of the current group.
# count(x) : counts number of repetitions of each element in a qualitative column x
# sum(!is.na(x)) : number of non-missing values in current group
# n_distinct(x) : number of distinct (unique) values in current group

# Counts and proportions of logical values: sum(x > 10), mean(y == 0)

# quantile(x, 0.25) will find a value of x that is greater than 25% of the values, and less than the rest
# IQR is 3rd Quartile - 1st Quartile (i.e the box plot)
# mad is median absolute deviation mad(x) may be more useful if you have outliers

library(nycflights13)
library(tidyverse)

(not_cancelled <- flights %>%
  filter(!is.na(arr_delay) & !is.na(dep_delay))
)

```

```

## # A tibble: 327,346 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     1     1     517           515         2      830           819
## 2  2013     1     1     533           529         4      850           830
## 3  2013     1     1     542           540         2      923           850
## 4  2013     1     1     544           545        -1     1004          1022
## 5  2013     1     1     554           600        -6      812           837
## 6  2013     1     1     554           558        -4      740           728
## 7  2013     1     1     555           600        -5      913           854
## 8  2013     1     1     557           600        -3      709           723
## 9  2013     1     1     557           600        -3      838           846
## 10 2013     1     1     558           600        -2      753           745
## # ... with 327,336 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>

```

```

(not_cancelled %>%
  group_by(year, month, day) %>%
  summarise(avg_arr_delay = mean(arr_delay), avg_pos_arr_delay = mean(arr_delay[arr_delay > 0]))
)

```

```

## # A tibble: 365 x 5
## # Groups:   year, month [12]
##   year month   day avg_arr_delay avg_pos_arr_delay
##   <int> <int> <int>         <dbl>         <dbl>
## 1  2013     1     1         12.7          32.5
## 2  2013     1     2         12.7          32.0
## 3  2013     1     3          5.73          27.7
## 4  2013     1     4        -1.93          28.3
## 5  2013     1     5        -1.53          22.6
## 6  2013     1     6          4.24          24.4
## 7  2013     1     7        -4.95          27.8
## 8  2013     1     8        -3.23          20.8
## 9  2013     1     9        -0.264          25.6
## 10 2013     1    10        -5.90          27.3

```

```
## # ... with 355 more rows
```

```
not_cancelled %>%  
  group_by(dest) %>%  
  summarise(distance_sd = sd(distance)) %>%  
  arrange(desc(distance_sd))
```

```
## # A tibble: 104 x 2  
##   dest distance_sd  
##   <chr>         <dbl>  
## 1 EGE          10.5  
## 2 SAN          10.4  
## 3 SFO          10.2  
## 4 HNL          10.0  
## 5 SEA           9.98  
## 6 LAS           9.91  
## 7 PDX           9.87  
## 8 PHX           9.86  
## 9 LAX           9.66  
## 10 IND          9.46
```

```
## # ... with 94 more rows
```

```
not_cancelled %>%  
  group_by(year, month, day) %>%  
  summarise(first = min(dep_time), last=max(dep_time))
```

```
## # A tibble: 365 x 5  
## # Groups:   year, month [12]  
##   year month day first last  
##   <int> <int> <int> <int> <int>  
## 1 2013     1     1   517 2356  
## 2 2013     1     2    42 2354  
## 3 2013     1     3    32 2349  
## 4 2013     1     4    25 2358  
## 5 2013     1     5    14 2357  
## 6 2013     1     6    16 2355  
## 7 2013     1     7    49 2359  
## 8 2013     1     8   454 2351  
## 9 2013     1     9     2 2252  
## 10 2013     1    10     3 2320
```

```
## # ... with 355 more rows
```

```
not_cancelled %>%  
  group_by(year, month, day) %>%  
  summarise(first_dep = first(dep_time), last_dep = last(dep_time))
```

```
## # A tibble: 365 x 5  
## # Groups:   year, month [12]  
##   year month day first_dep last_dep  
##   <int> <int> <int>     <int>     <int>  
## 1 2013     1     1       517       2356  
## 2 2013     1     2        42       2354  
## 3 2013     1     3        32       2349  
## 4 2013     1     4        25       2358  
## 5 2013     1     5        14       2357  
## 6 2013     1     6        16       2355
```

```
## 7 2013 1 7 49 2359
## 8 2013 1 8 454 2351
## 9 2013 1 9 2 2252
## 10 2013 1 10 3 2320
## # ... with 355 more rows
```

Filtering on ranks gives you all variables, with each observation in a separate row:

```
not_cancelled %>%
  group_by(year, month, day) %>%
  mutate(rank = min_rank(desc(dep_time))) %>%
  filter(rank %in% range(rank))
```

```
## # A tibble: 770 x 20
```

```
## # Groups:   year, month, day [365]
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1 2013     1     1     517           515         2      830           819
## 2 2013     1     1    2356           2359        -3      425           437
## 3 2013     1     2      42           2359        43      518           442
## 4 2013     1     2    2354           2359        -5      413           437
## 5 2013     1     3      32           2359        33      504           442
## 6 2013     1     3    2349           2359       -10      434           445
## 7 2013     1     4      25           2359        26      505           442
## 8 2013     1     4    2358           2359        -1      429           437
## 9 2013     1     4    2358           2359        -1      436           445
## 10 2013     1     5     14           2359        15      503           445
```

```
## # ... with 760 more rows, and 12 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>,
## #   rank <int>
```

which destination have the most carriers

```
not_cancelled %>%
  filter(!is.na(dest) & !is.na(carrier)) %>%
  group_by(dest) %>%
  summarize(max_carriers = max(n_distinct(carrier))) %>%
  arrange(desc(max_carriers))
```

```
## # A tibble: 104 x 2
```

```
##   dest max_carriers
##   <chr>         <int>
## 1 ATL             7
## 2 BOS             7
## 3 CLT             7
## 4 ORD             7
## 5 TPA             7
## 6 AUS             6
## 7 DCA             6
## 8 DTW             6
## 9 IAD             6
## 10 MSP            6
```

```
## # ... with 94 more rows
```

Give a count of each destinations separately

```
not_cancelled %>%
  count(dest)
```

```
## # A tibble: 104 x 2
##   dest      n
##   <chr> <int>
## 1 ABQ     254
## 2 ACK     264
## 3 ALB     418
## 4 ANC       8
## 5 ATL   16837
## 6 AUS     2411
## 7 AVL     261
## 8 BDL     412
## 9 BGR     358
## 10 BHM     269
## # ... with 94 more rows
```

```
# You can optionally provide a weight variable.
# "count" (sum) the total number of miles a plane flew:
```

```
not_cancelled %>%
  count(tailnum, wt=distance)
```

```
## # A tibble: 4,037 x 2
##   tailnum      n
##   <chr>   <dbl>
## 1 D942DN    3418
## 2 NOEGMQ  239143
## 3 N10156 1096664
## 4 N102UW   25722
## 5 N103US   24619
## 6 N104UW   24616
## 7 N10575 139903
## 8 N105UW   23618
## 9 N107US   21677
## 10 N108UW   32070
## # ... with 4,027 more rows
```

```
# Which is the same as
not_cancelled %>%
  group_by(tailnum) %>%
  summarise(n=sum(distance))
```

```
## # A tibble: 4,037 x 2
##   tailnum      n
##   <chr>   <dbl>
## 1 D942DN    3418
## 2 NOEGMQ  239143
## 3 N10156 1096664
## 4 N102UW   25722
## 5 N103US   24619
## 6 N104UW   24616
## 7 N10575 139903
## 8 N105UW   23618
## 9 N107US   21677
## 10 N108UW   32070
## # ... with 4,027 more rows
```

```
# How many flights left before 5am?
not_cancelled %>%
  group_by(year, month, day) %>%
  summarise(n_early = sum(dep_time < 500))
```

```
## # A tibble: 365 x 4
## # Groups:   year, month [12]
##   year month   day n_early
##   <int> <int> <int>   <int>
## 1  2013     1     1         0
## 2  2013     1     2         3
## 3  2013     1     3         4
## 4  2013     1     4         3
## 5  2013     1     5         3
## 6  2013     1     6         2
## 7  2013     1     7         2
## 8  2013     1     8         1
## 9  2013     1     9         3
## 10 2013     1    10         3
## # ... with 355 more rows
```

```
# What proportion of flights are delayed by more than an hour?
not_cancelled %>%
  group_by(year, month, day) %>%
  summarize(proportion = mean(arr_delay > 60))
```

```
## # A tibble: 365 x 4
## # Groups:   year, month [12]
##   year month   day proportion
##   <int> <int> <int>       <dbl>
## 1  2013     1     1    0.0722
## 2  2013     1     2    0.0851
## 3  2013     1     3    0.0567
## 4  2013     1     4    0.0396
## 5  2013     1     5    0.0349
## 6  2013     1     6    0.0470
## 7  2013     1     7    0.0333
## 8  2013     1     8    0.0213
## 9  2013     1     9    0.0202
## 10 2013     1    10    0.0183
## # ... with 355 more rows
```

```
library(nycflights13)
library(tidyverse)

(per_day <- flights %>%
  group_by(year, month, day) %>%
  summarise(flights=n())
)
```

```
## # A tibble: 365 x 4
## # Groups:   year, month [12]
##   year month   day flights
##   <int> <int> <int>   <int>
## 1  2013     1     1     842
```



```
## 2 2013 1 2 943
## 3 2013 1 3 914
## 4 2013 1 4 915
## 5 2013 1 5 720
## 6 2013 1 6 832
## 7 2013 1 7 933
## 8 2013 1 8 899
## 9 2013 1 9 902
## 10 2013 1 10 932
## # ... with 355 more rows
```

```
(per_month <- per_day %>%
  summarize(flights = sum(flights)))
```

```
## # A tibble: 12 x 3
## # Groups:   year [1]
##   year month flights
##   <int> <int>   <int>
## 1 2013     1  27004
## 2 2013     2  24951
## 3 2013     3  28834
## 4 2013     4  28330
## 5 2013     5  28796
## 6 2013     6  28243
## 7 2013     7  29425
## 8 2013     8  29327
## 9 2013     9  27574
## 10 2013    10  28889
## 11 2013    11  27268
## 12 2013    12  28135
```

```
(per_year <- per_month %>%
  summarize(flights = sum(flights)))
```

```
## # A tibble: 1 x 2
##   year flights
##   <int>   <int>
## 1 2013  336776
```

Equivalently

```
(per_day <- flights %>%
  group_by(year, month, day) %>%
  summarise(per_day_flights=n()) %>%
  summarise(per_month_flights = sum(per_day_flights)) %>%
  summarise(per_year_flights = sum(per_month_flights))
)
```

```
## # A tibble: 1 x 2
##   year per_year_flights
##   <int>           <int>
## 1 2013           336776
```

Be careful when progressively rolling up summaries: it's OK for sums and counts, but you need to think about medians.
I.e. the sum of groupwise sums is the overall sum, but the median of groupwise medians is not the overall median.

If you need to remove grouping, and return to operations on ungrouped data, use ungroup()

```

daily <- group_by(flights, year, month, day)

daily %>%
  ungroup() %>% # no longer grouped by year-month-day
  summarise(flights = n())

## # A tibble: 1 x 1
##   flights
##   <int>
## 1 336776

library(nycflights13)
library(tidyverse)

not_cancelled <- flights %>%
  filter(!is.na(arr_delay) & (!is.na(dep_delay)))
# 1)
# A flight is 15 minutes early 50% of the time, and 15 minutes late 50% of the time.
(not_cancelled %>%
  group_by(flight) %>%
  summarize(total = n(),
            early15 = sum(arr_delay == -15),
            late15 = sum(arr_delay == 15)) %>%
  filter(total != 0 & early15 != 0 & late15 != 0 & near(total / 2, 0.5))
)

## # A tibble: 0 x 4
## # ... with 4 variables: flight <int>, total <int>, early15 <int>, late15 <int>

# Another way :
(not_cancelled %>%
  group_by(flight) %>%
  summarize(total = n(),
            early15 = mean(arr_delay == -15, na.rm = T),
            late15 = mean(arr_delay == 15, na.rm = T)) %>%
  filter(total != 0 & early15 == 0.5 & late15 == 0.5)
)

## # A tibble: 0 x 4
## # ... with 4 variables: flight <int>, total <int>, early15 <dbl>, late15 <dbl>

# A flight is always 10 minutes late.
(not_cancelled %>%
  group_by(flight) %>%
  filter(arr_delay == 10)
)

## # A tibble: 3,373 x 19
## # Groups:   flight [1,475]
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1 2013     1     1     624           630          -6     840           830
## 2 2013     1     1     717           720          -3     850           840
## 3 2013     1     1     745           745           0    1135          1125

```

```
## 4 2013 1 1 805 805 0 1015 1005
## 5 2013 1 1 811 815 -4 1026 1016
## 6 2013 1 1 921 900 21 1237 1227
## 7 2013 1 1 1158 1205 -7 1530 1520
## 8 2013 1 1 1211 1215 -4 1423 1413
## 9 2013 1 1 1455 1459 -4 1655 1645
## 10 2013 1 1 1554 1600 -6 1830 1820
## # ... with 3,363 more rows, and 11 more variables: arr_delay <dbl>,
## # carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## # air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# 99% of the time a flight is on time. 1% of the time it's 2 hours late.
```

```
not_cancelled %>%
  group_by(flight) %>%
  summarize (total = n(), ontime = sum(arr_delay == 0), late = sum(arr_delay == 2)) %>%
  filter((ontime %/% total)*100 == 99 && (late %/% total)*100 == 1)
```

```
## # A tibble: 0 x 4
## # ... with 4 variables: flight <int>, total <int>, ontime <int>, late <int>
```

```
# 2)
# not_cancelled %>% count(dest)
not_cancelled %>%
  group_by(dest) %>%
  summarise(n = n())
```

```
## # A tibble: 104 x 2
##   dest      n
##   <chr> <int>
## 1 ABQ    254
## 2 ACK    264
## 3 ALB    418
## 4 ANC      8
## 5 ATL  16837
## 6 AUS   2411
## 7 AVL    261
## 8 BDL    412
## 9 BGR    358
## 10 BHM   269
## # ... with 94 more rows
```

```
# not_cancelled %>% count(tailnum, wt = distance)
not_cancelled %>%
  group_by(tailnum) %>%
  summarise(wt = sum(distance))
```


```
## # A tibble: 4,037 x 2
##   tailnum      wt
##   <chr>    <dbl>
## 1 D942DN    3418
## 2 NOEGMQ  239143
## 3 N10156  109664
## 4 N102UW   25722
## 5 N103US   24619
## 6 N104UW   24616
```

```
## 7 N10575 139903
## 8 N105UW 23618
## 9 N107US 21677
## 10 N108UW 32070
## # ... with 4,027 more rows
```

```
# 4)
# Look at the number of cancelled flights per day. Is there a pattern? Is the proportion of cancelled f
```

```
flights %>%
  group_by(year, month, day) %>%
  summarize(cancelledFlights = sum(is.na(arr_delay) | (is.na(dep_delay))), avgArrDelay = mean(is.na(arr_
  arrange(desc(avgArrDelay)) %>%
  filter(cancelledFlights <= 75) %>%
  ggplot(mapping = aes(x=cancelledFlights, y = avgArrDelay)) +
  geom_point()+
  geom_smooth(se=F)
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
# 5)
# Which carrier has the worst delays?
# Challenge: can you disentangle the effects of bad airports vs. bad carriers? Why/why not?
# (Hint: think about flights %>% group_by(carrier, dest) %>% summarise(n()))
```

```
# not_cancelled %>%
# group_by(carrier, dest) %>%
```

```
# summarise(max_delay_per_dest = max(arr_delay, na.rm = T)) %>%
# summarise(max_delay_per_carr = max(max_delay_per_dest, na.rm = T)) %>%
# arrange(desc(max_delay_per_carr))
```

```
not_cancelled %>%
  group_by(carrier, dest) %>%
  summarise(max_delay_per_dest = max(arr_delay, na.rm = T)) %>%
  group_by(dest) %>%
  mutate(rank = min_rank(desc(max_delay_per_dest))) %>%
  filter(rank %in% range(rank)) %>%
  arrange(carrier, dest)
```

```
## # A tibble: 180 x 4
## # Groups:   dest [104]
##   carrier dest max_delay_per_dest rank
##   <chr>   <chr>          <dbl> <int>
## 1 9E     ATL             55      7
## 2 9E     AUS             25      6
## 3 9E     AVL             13      2
## 4 9E     BTV             -1      3
## 5 9E     BUF            396      1
## 6 9E     CAE             55      2
## 7 9E     CLT            744      1
## 8 9E     CMH             70      3
## 9 9E     DAY            292      1
## 10 9E    DCA            384      1
## # ... with 170 more rows
```

```
# 6) What does the sort argument to count() do?
not_cancelled %>%
  count(tailnum, wt = distance, sort = T)
```

```
## # A tibble: 4,037 x 2
##   tailnum      n
##   <chr>   <dbl>
## 1 N328AA 929090
## 2 N338AA 921172
## 3 N335AA 902271
## 4 N327AA 900482
## 5 N323AA 839468
## 6 N319AA 837924
## 7 N336AA 833136
## 8 N329AA 825826
## 9 N324AA 786159
## 10 N339AA 783648
## # ... with 4,027 more rows
```

```
# assume we want to count (number, letter) pair
(data = tibble(
  letter = sample(LETTERS, 50000, replace = TRUE),
  number = sample(1:10, 50000, replace = TRUE)
))
```

```
## # A tibble: 50,000 x 2
##   letter number
```

```
##      <chr>      <int>
## 1 D              4
## 2 R             10
## 3 E             10
## 4 R              6
## 5 M              3
## 6 G              4
## 7 F              8
## 8 A              1
## 9 J              2
## 10 U             9
## # ... with 49,990 more rows
```

```
data %>%
  count(letter, number, sort = TRUE)
```

```
## # A tibble: 260 x 3
##   letter number     n
##   <chr>    <int> <int>
## 1 E              8  233
## 2 F              5  229
## 3 T              9  227
## 4 Q              9  224
## 5 J              4  222
## 6 I              4  221
## 7 P              8  221
## 8 F              2  220
## 9 R              3  220
## 10 S             7  220
## # ... with 250 more rows
```

```
data %>%
  group_by(letter, number) %>%
  summarise(n = n()) %>%
  ungroup() %>%
  arrange(desc(n))
```

```
## # A tibble: 260 x 3
##   letter number     n
##   <chr>    <int> <int>
## 1 E              8  233
## 2 F              5  229
## 3 T              9  227
## 4 Q              9  224
## 5 J              4  222
## 6 I              4  221
## 7 P              8  221
## 8 F              2  220
## 9 R              3  220
## 10 S             7  220
## # ... with 250 more rows
```

```
data %>%
  count(letter, number) %>%
  ungroup() %>%
  arrange(desc(n))
```

```
## # A tibble: 260 x 3
##   letter number     n
##   <chr>    <int> <int>
## 1 E             8   233
## 2 F             5   229
## 3 T             9   227
## 4 Q             9   224
## 5 J             4   222
## 6 I             4   221
## 7 P             8   221
## 8 F             2   220
## 9 R             3   220
## 10 S            7   220
## # ... with 250 more rows
```

```
library(nycflights13)
library(tidyverse)
```

```
# Find the worst members of each group
flights %>%
  group_by(year, month, day) %>%
  filter(rank(desc(arr_delay)) < 4)
```

```
## # A tibble: 1,105 x 19
## # Groups:   year, month, day [365]
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>    <int>         <int>
## 1 2013     1     1     848           1835           853     1001           1950
## 2 2013     1     1    1815           1325           290     2120           1542
## 3 2013     1     1    2343           1724           379      314           1938
## 4 2013     1     2    1412           838           334     1710           1147
## 5 2013     1     2    1607          1030           337     2003           1355
## 6 2013     1     2    2131          1512           379     2340           1741
## 7 2013     1     3    2008          1540           268     2339           1909
## 8 2013     1     3    2012          1600           252     2314           1857
## 9 2013     1     3    2056          1605           291     2239           1754
## 10 2013     1     4    1305          1030           155     1452           1210
## # ... with 1,095 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
(r1 <- rank(x1 <- c(3, 1, 4, 15, 92)))
```

```
## [1] 2 1 3 4 5
```

```
(r2 <- min_rank(x1 <- c(3, 1, 4, 15, 92)))
```

```
## [1] 2 1 3 4 5
```

```
# Find all groups bigger than a threshold
(pouplar_dests <-
  flights %>%
  group_by(dest) %>%
  filter(n() > 365))
```

```
## # A tibble: 332,577 x 19
## # Groups:   dest [77]
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
## 1  2013     1     1     517             515         2     830             819
## 2  2013     1     1     533             529         4     850             830
## 3  2013     1     1     542             540         2     923             850
## 4  2013     1     1     544             545        -1    1004            1022
## 5  2013     1     1     554             600        -6     812             837
## 6  2013     1     1     554             558        -4     740             728
## 7  2013     1     1     555             600        -5     913             854
## 8  2013     1     1     557             600        -3     709             723
## 9  2013     1     1     557             600        -3     838             846
## 10 2013     1     1     558             600        -2     753             745
## # ... with 332,567 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

A grouped filter is a grouped mutate followed by an ungrouped filter. I generally avoid them except for

```
pouplar_dests %>%
  filter(arr_delay > 0) %>%
  mutate(prop_delay = arr_delay / sum(arr_delay)) %>%
  select(year:day, dest, arr_delay, prop_delay)
```

```
## # A tibble: 131,106 x 6
## # Groups:   dest [77]
##   year month   day dest  arr_delay prop_delay
##   <int> <int> <int> <chr>    <dbl>      <dbl>
## 1  2013     1     1 IAH      11  0.000111
## 2  2013     1     1 IAH      20  0.000201
## 3  2013     1     1 MIA      33  0.000235
## 4  2013     1     1 ORD      12  0.0000424
## 5  2013     1     1 FLL      19  0.0000938
## 6  2013     1     1 ORD       8  0.0000283
## 7  2013     1     1 LAX       7  0.0000344
## 8  2013     1     1 DFW      31  0.000282
## 9  2013     1     1 ATL      12  0.0000400
## 10 2013     1     1 DTW      16  0.000116
## # ... with 131,096 more rows
```

```
# 1)
# Filter function is applied to each group and shrinks the elements of each group
flights %>%
  group_by(year, month, day) %>%
  filter (air_time == 320 & carrier=="US")
```

```
## # A tibble: 15 x 19
## # Groups:   year, month, day [15]
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
## 1  2013     1    11    1619             1620        -1    2003             2003
## 2  2013     1    14    1346             1350        -4    1724             1715
## 3  2013    10    30     625             630        -5     931             918
## 4  2013    11    17    1010             1015        -5    1406             1342
## 5  2013    11    22    1639             1630         9    2022             2011
## 6  2013    11    23    1030             1030         0    1417             1410
```



```
## 7 2013 12 4 1352 1355 -3 1729 1715
## 8 2013 12 5 1402 1355 7 1736 1715
## 9 2013 12 6 1406 1355 11 1746 1715
## 10 2013 12 20 1631 1630 1 2009 1957
## 11 2013 2 23 954 959 -5 1345 1333
## 12 2013 3 23 1007 1015 -8 1252 1240
## 13 2013 4 1 623 630 -7 922 913
## 14 2013 4 2 1628 1635 -7 1902 1856
## 15 2013 4 26 1632 1630 2 1910 1851
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# mutate function with group
```

```
# Arithmetic operators with group_by
```

```
# flights %>%
```

```
#Functions that work most naturally in grouped mutates and filters are known as window functions (vs. t
```

```
vignette("window-functions")
```

```
## starting httpd help server ... done
```